B.Comp. Dissertation CA Report

Benchmarking and Improving OCR System for Southeast Asian Languages

By

Qiu Jiasheng, Jason

Department of Computer Science

School of Computing

National University of Singapore

2024/2025

B.Comp. Dissertation CA Report

Benchmarking and Improving OCR System for Southeast Asian Languages

By

Qiu Jiasheng, Jason

Department of Computer Science

School of Computing

National University of Singapore

2024/2025

Project ID: H0792230

Supervisor: A/P Min-Yen Kan

Advisor: Tongyao Zhu

Abstract

While Optical Character Recognition (OCR) has been widely studied for high-resource languages such as English and Chinese, the efficacy and limitations of OCR models on Southeast Asian (SEA) languages remain largely unexplored. This study aims to bridge this gap by assessing and improving the performance of OCR technologies on SEA languages. To achieve this objective, we propose a reusable pipeline to gather SEA-language text from Wikipedia and benchmark popular OCR tools.

Subject Descriptors:

I.2.7 Natural Language Processing

Keywords:

Optical Character Recognition, Southeast Asian Languages

Implementation Software and Hardware:

Python, Tesseract, EasyOCR

Acknowledgement

I would like to thank my supervisor, A/P Kan Min-Yen, and my advisor, Tongyao Zhu, for their invaluable guidance and mentorship. Their encouragement and constructive guidance have been a significant source of inspiration throughout the project.

List of Figures

List of Tables

Table of Contents

Chapter 1

Introduction

Current research in Natural Language Processing (NLP) is heavily concentrated on 20 of the 7,000 languages in the world (Magueresse et al., 2020). In particular, Southeast Asia (SEA) is home to over 1,000 languages but remains a relatively under-researched region in NLP (Aji et al., 2023). Similar to most low-resource languages, a major challenge in developing NLP systems for SEA languages is the limited availability of datasets for the region's languages. Although many scanned documents and books in these low-resource languages are available online, the text within these files remains inaccessible due to formats like images and PDFs.

A solution to this problem is to use Optical Character Recognition (OCR) to extract the textual data. OCR is the process of identifying and converting text in an image into a computer-friendly text format. By extracting the text from these scanned documents, OCR can generate valuable datasets for low-resource languages. The created datasets can then be used for downstream NLP tasks, such as machine translation, training large language models, and named-entity recognition (Agarwal & Anastasopoulos, 2024; Ignat et al., 2022). Therefore, studying OCR performance on SEA languages is crucial to accelerating NLP research in the region.

While OCR has been widely studied for high-resource languages such as English and Chinese, the efficacy and limitations of OCR models on SEA languages

remain largely unexplored. To address this gap, we propose a reusable pipeline to collect textual data in low-resource SEA languages from Wikipedia and benchmark popular open-source OCR tools on the collected data. The primary objective is to benchmark and improve the performance of OCR technologies on SEA languages, thereby contributing to the advancement of NLP applications in this linguistically diverse region. Specifically, this project seeks to answer the following research questions (RQs):

- **RQ1.** How do popular OCR tools perform on SEA scripts?

- **RQ2.** What specific linguistic and script-related challenges affect OCR accuracy on SEA languages?

- **RQ3.** What techniques and recommendations can enhance OCR accuracy on SEA languages?

Chapter 2

Related Work

2.1 Collecting Low-Resource Language Data

To evaluate OCR performance accurately, a collection of textual data in the form of images or PDFs paired with reliable ground truth is needed. Similar to most NLP tasks, data scarcity poses a major obstacle to advancing OCR technology in low-resource languages, where the limited availability of annotated textual data restricts both model training and evaluation. Although an abundance of scanned documents in these low-resource languages exists online, they lack the ground truth required for evaluation. While plain text in these languages is often available separately, it typically exists in text-based formats rather than images or PDFs, limiting its direct usefulness as ground truth for OCR.

To bridge this gap, many studies rely on artificial images and PDFs generated from plain text to create usable evaluation data. For instance, Ignat et al. (2022) artificially created PDFs from the Flores 101 dataset, which consists of text data from Wikipedia in 101 languages. Generalizing this concept of a document creation pipeline further, Gupte et al. (2021) published an open-source Python package for generating document images from plain text, including several document styling templates. Using these methods, high-quality low-resource language data paired with ground truth can be generated from text-based formats on a large scale.

A common trend in using artificial data is the augmentation of noise to simulate real-world conditions, which often contain complex layouts, stains, and scribbles (Hegghammer, 2022). Directly using noise-free data, i.e., single-column text in a clear font, limits OCR processors' usefulness on real-life scanned documents. Thus, noise augmentation is often applied to artificial data. Some popular techniques include changing font style, size, color, letter spacing, and adding Gaussian blur, bleed-through, and salt-and-pepper noise (Gupte et al., 2021; Ignat et al., 2022).

2.2   Existing OCR Systems

To benchmark OCR performance programmatically, there are two broad categories of OCR systems: open-source and commercial.

Open-source OCR systems are characterized by their accessibility, allowing users to view, modify, and distribute the source code freely. This transparency enables developers to customize the tools to meet specific needs. Furthermore, open-source systems typically incur no licensing fees, making them cost-effective options for research purposes. For instance, Tesseract (Smith, 2013) is a popular open-source OCR engine mentioned consistently in related studies (Hegghammer, 2022; Ignat et al., 2022).

Commercial OCR systems are typically accessed through paid services via application programming interfaces (APIs). Notable examples of these systems include Amazon Textract, Google Document AI, and Google Vision API. Generally, commercial off-the-shelf OCR tools tend to perform better than their open-source counterparts (Hegghammer, 2022; Ignat et al., 2022). However, the proprietary

nature of these "black box" models limits their utility for research purposes, such as fine-tuning and customization. Additionally, the associated costs of these services contribute to fewer studies focused on commercial systems in comparison to the more widely researched open-source tools.

2.3   Benchmarking OCR

Applying OCR yields a plain text prediction, which is then compared with the ground truth data to assess the tool's accuracy and performance. Recent studies have demonstrated that OCR systems tend to perform better on artificially generated data than on real-world data (Ignat et al., 2022). This observation suggests that synthetic datasets may not fully capture the complexities of authentic documents, which often feature issues like imperfect text alignment, varied fonts, and complex layouts. Furthermore, the addition of synthetic noise significantly raises error rates, especially for open-source systems, which appear more susceptible to noise interference than their commercial counterparts (Hegghammer, 2022).

When comparing performance on different scripts, OCR tools generally achieve higher accuracy on scripts written in Latin alphabets (Hegghammer, 2022; Ignat et al., 2022). This disparity in performance partly stems from market incentives that prioritize the development of English-language OCR systems, resulting in more extensive training data and refinement for Latin-based scripts. Ornate scripts, such as those with complex diacritics or unique letter shapes, present additional challenges and tend to yield lower OCR accuracy.

In terms of benchmarking OCR on SEA languages, the most related study is

the recent work by Ignat et al. (2022). They grouped 60 low-resource languages by region and script, including SEA languages like Khmer, Lao, Burmese, Thai, and Vietnamese. Their research revealed that while OCR models perform well on artificial SEA-language data, accuracy drops significantly on real-world data. This discrepancy underscores the need for more real-world training data to improve OCR outcomes for SEA languages.

In summary, there exists a gap in benchmarking OCR specifically for SEA languages, largely due to the lack of real-world training data. This project addresses this gap by proposing a reusable pipeline for benchmarking OCR performance on real-world SEA-language data sourced from Wikipedia.

Chapter 3

Methodology

3.1   Data Selection

Source of Data

We chose to use Wikipedia as our text corpus for several reasons. Firstly, Wikipedia articles can be easily converted into images via screenshots, making them suitable for OCR applications. The platform also offers a convenient source of ground truth through its APIs that provide plain text for most articles. Secondly, Wikipedia hosts a large corpus in several popular SEA languages, including Thai, Vietnamese, Indonesian, Tamil, and Burmese, supporting our language needs ("List of Wikipedias", 2024). Lastly, Wikipedia articles contain visual elements like images and tables that are common in modern real-world documents.

Languages

From the languages available on Wikipedia, we selected English, Indonesian, Vietnamese, and Thai text. English serves as a baseline for sanity checks and bug fixing. The remaining SEA languages were chosen to capture diverse script characteristics. Indonesian represents Latin-based scripts, Vietnamese represents Latin scripts with diacritics, and Thai represents non-Latin scripts.

Dataset

We compiled a dataset of 100 Wikipedia articles, selecting the 20 most viewed English articles from each of five categories: people, present countries, cities[1], life, and buildings and structures[2] ("Wikipedia:Popular pages", 2024). These categories were chosen to create a diverse corpus in terms of content. We prioritized popular articles because they are more likely to have substantial data available across languages. Table A.1 lists the articles included in our dataset.
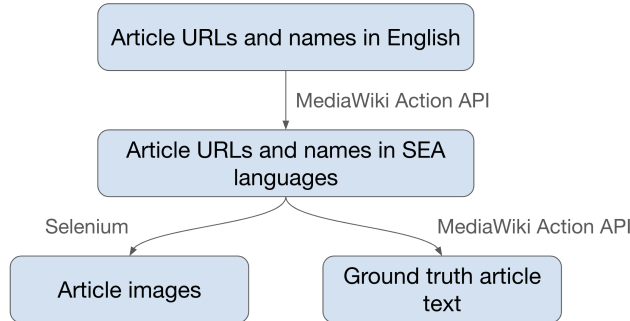
3.2 Data Collection

Figure 3.1: Pipeline for data collection from Wikipedia

From the dataset of 100 Wikipedia articles, we collected article images and ground truth article text in our selected languages using Python, Selenium[3], and the MediaWiki Action API[4]. Figure 3.1 illustrates the overall pipeline for data collection. The detailed steps are as follows:

---

[1]Singapore was replaced because it's already listed under present countries.

[2]Machu Picchu was replaced because it's already listed under cities.

[3]Selenium is a framework for automating web browsers, commonly used for web scraping by programmatically interacting with websites.

[4]The MediaWiki Action API allows access to wiki page operation features such as search and retrieval.

1. Manually compile the dataset's article names and URLs in English.

2. Fetch the article names and URLs in Thai, Vietnamese, and Indonesian from the MediaWiki Action API.

3. Download the article PDFs in all languages using Selenium.

4. Convert the article PDFs into PNG images, where each image represents one page in the PDF.

5. Download the ground truth article text into TXT files from the MediaWiki Action API.

3.3 OCR Evaluation

OCR Systems

In our selection of OCR systems for benchmarking, we prioritized open-source solutions that support a diverse range of SEA languages, as this approach enhances accessibility and reusability for the proposed evaluation pipeline. Consequently, we selected to use Tesseract and EasyOCR.

Tesseract[5] is an established OCR engine, recognized as one of the top performers in the 1995 UNLV Test (Rice et al., 1995). It utilizes an underlying Long Short-Term Memory (LSTM) model. EasyOCR[6] is a modern OCR framework that integrates a text detection model based on the Character Region Awareness for Text (CRAFT) algorithm with a recognition model utilizing a Convolutional

---

[5]https://github.com/tesseract-ocr/tesseract
[6]https://github.com/JaidedAI/EasyOCR

Recurrent Neural Network (CRNN). Both Tesseract and EasyOCR provide robust support for English, Indonesian, Vietnamese, and Thai, making them suitable candidates for our benchmarking study.

Evaluation Metrics

$$CER = \frac{I + D + S}{N} \tag{3.1}$$

Similar to most OCR benchmark studies, we utilize Character Error Rate (CER) and Word Error Rate (WER) as our evaluation metrics (Hegghammer, 2022; Ignat et al., 2022). CER measures the accuracy of character recognition and is calculated using the Levenshtein or edit distance, which represents the minimum number of single-character insertions (I), deletions (D), and substitutions (S) required to transform one word into another. As shown in Equation 3.1, CER is defined as the edit distance between the OCR-predicted text and ground truth text, divided by the total number of characters in the ground truth text (N). A lower CER value indicates higher accuracy, with 0 representing perfect recognition. Notably, CER can exceed 1, particularly when there are a significant number of insertions. WER serves as the word-based counterpart to CER.

Data Validation and Cleaning

The raw predicted text generated by the OCR tools exhibited extremely high error rates. Upon further investigation, we identified some consistent formatting issues in the output:

- Tesseract adds an additional space character after every predicted character, leading to CERs exceeding 1.

- The article images included references and in-text citations, which are not present in the ground truth.

To address these issues, we performed data cleaning to align the output text more closely with the ground truth. This cleaning significantly reduced the error rates to more acceptable levels.

We then implemented a data validation step to automatically check the CERs for each language. Articles were flagged as outliers if the CER between the OCR-predicted text and the ground truth text exceeded two standard deviations from the mean (Cousineau & Chartier, 2010). We manually reviewed these outlier articles for anomalies, resulting in the removal of seven articles where the images and ground truth texts contained different content.

Summary

In summary, we evaluated the CER and WER for each article using the images and ground truth text collected from Wikipedia. Figure 3.2 illustrates the overall pipeline for OCR evaluation. The detailed steps are as follows:

1. Apply the OCR tools on the article images.

2. Perform data cleaning on the predicted article text and ground truth article text.
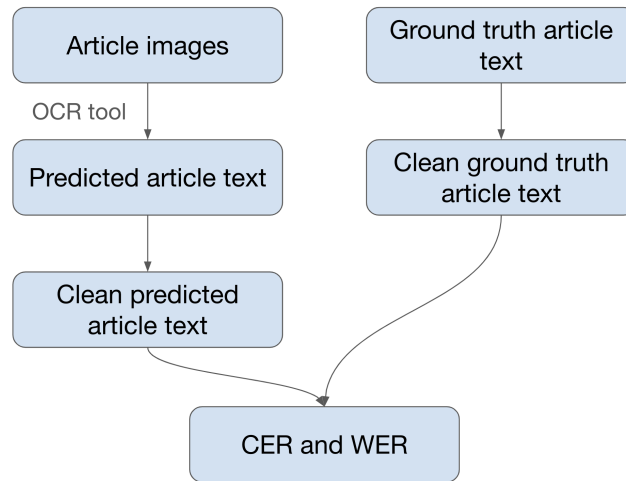
Figure 3.2: Pipeline for OCR evaluation

3. Compute the CER and WER between the predicted text and the ground truth text using JiWER[7].

---

[7]JiWER is a Python package designed for fast calculation of CER and WER.

Chapter 4

Results

	Character Error Rate		Word Error Rate	
	EasyOCR	Tesseract	EasyOCR	Tesseract
English	0.17	0.20	0.25	0.29
Indonesian	0.20	0.18	0.27	0.33
Vietnamese	0.30	0.39	0.31	0.42
Thai	0.26	0.51	1.68	1.77

Table 4.1: Average Character Error Rate and Word Error Rate

Table 4.1 presents the average CER and WER for each language and OCR model. The results reveal several key observations on the performance of popular OCR tools on SEA languages (**RQ1**):

- EasyOCR generally outperforms Tesseract, yielding lower error rates.

- Both OCR tools perform best on Latin scripts (English and Indonesian), followed by Latin scripts with diacritics (Vietnamese), and then Non-Latin scripts (Thai).

- The WER for Thai is notably high due to its non-segmented script structure.

4.1   Discussion

There is no universal standard for what makes a "good" CER or WER value, as OCR performance varies widely depending on factors like text complexity, layout,

and print quality. While many OCR systems claim accuracy rates of 99%, Holley (2009) argues that these rates typically apply to clean images or to cases with manual intervention during the OCR process. Supporting this claim, Hegghammer (2022) found that his WER on English articles averaged around 0.01 (nearly perfect) on clean images but surged to a WER of 0.22 when noise was introduced. For comparison, our average WER on English articles is 0.27. Therefore, our results align with similar studies that benchmark OCR on noisy, real-world data.
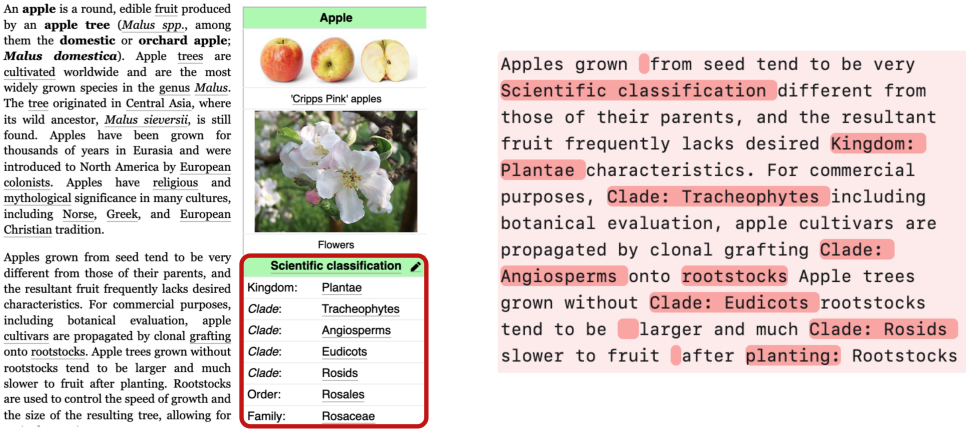

Figure 4.1: The OCR tools detect visual elements, like the table in the left figure, as additional text highlighted in the right figure.

Manual comparison of the OCR-predicted text and the ground truth reveals that most errors arise from Wikipedia's layout and graphical content, such as tables and embedded text within images. As shown in Figure 4.1, OCR tools often detect these elements as additional text, which is absent from the ground truth, leading to discrepancies. This issue highlights the challenge OCR systems face with complex layouts and visual elements that are misinterpreted as part of the textual content.

Chapter 5

Future Work

In summary, this study examines the performance of OCR systems on SEA languages. We developed a reusable pipeline for collecting real-world data from Wikipedia to generate images and ground truth text from 100 articles in English, Indonesian, Vietnamese, and Thai. Using this dataset, we benchmarked Tesseract and EasyOCR, yielding results comparable to those in recent studies.

For future studies, we hope to focus on the following areas:

- Expand our text corpus to other SEA languages.

- Try other more OCR tools (e.g., transformer-based)

- Identify language-specific challenges faced by OCR in the results (**RQ2**).

- Suggest techniques and recommendations to enhance OCR accuracy on SEA languages (**RQ3**).

References

Agarwal, M., & Anastasopoulos, A. (2024). A concise survey of OCR for low-resource languages. In M. Mager, A. Ebrahimi, S. Rijhwani, A. Oncevay, L. Chiruzzo, R. Pugh, & K. von der Wense (Eds.), *Proceedings of the 4th workshop on natural language processing for indigenous languages of the americas (americasnlp 2024)* (pp. 88–102). Association for Computational Linguistics. https://doi.org/10.18653/v1/2024.americasnlp-1.10

Aji, A. F., Forde, J. Z., Loo, A. M., Sutawika, L., Wang, S., Winata, G. I., Yong, Z.-X., Zhang, R., Doğruöz, A. S., Tan, Y. L., & Cruz, J. C. B. (2023). Current status of NLP in South East Asia with insights from multilingualism and language diversity. *Proceedings of the 13th International Joint Conference on Natural Language Processing and the 3rd Conference of the Asia-Pacific Chapter of the Association for Computational Linguistics*, 8–13. https://aclanthology.org/2023.ijcnlp-tutorials.2

Cousineau, D., & Chartier, S. (2010). Outliers detection and treatment: A review. *International Journal of Psychological Research*, *3*, 58–67. https://doi.org/10.21500/20112084.844

Gupte, A., Romanov, A., Mantravadi, S., Banda, D., Liu, J., Khan, R., Meenal, L. R., Han, B., & Srinivasan, S. (2021). Lights, camera, action! A framework to improve NLP accuracy over OCR documents. *CoRR*, *abs/2108.02899*. https://arxiv.org/abs/2108.02899

Hegghammer, T. (2022). OCR with Tesseract, Amazon Textract, and Google Document AI: A benchmarking experiment. *Journal of Computational Social*

Science, *5*, 861–882. https://doi.org/https://doi.org/10.1007/s42001-021-00149-1

Holley, R. (2009). How good can it get? analysing and improving ocr accuracy in large scale historic newspaper digitisation programs. *D-Lib Magazine*, *15*(3/4). https://doi.org/10.1045/march2009-holley

Ignat, O., Maillard, J., Chaudhary, V., & Guzmán, F. (2022). OCR improves machine translation for low-resource languages. In S. Muresan, P. Nakov, & A. Villavicencio (Eds.), *Findings of the association for computational linguistics: Acl 2022* (pp. 1164–1174). Association for Computational Linguistics. https://doi.org/10.18653/v1/2022.findings-acl.92

List of wikipedias. (2024). https://en.wikipedia.org/wiki/List_of_Wikipedias

Magueresse, A., Carles, V., & Heetderks, E. (2020). Low-resource languages: A review of past work and future challenges. *CoRR*, *abs/2006.07264*. https://arxiv.org/abs/2006.07264

Rice, S., Jenkins, F., & Nartker, T. (1995). *The fourth annual test of OCR accuracy* (tech. rep.). Information Science Research Institute.

Smith, R. W. (2013). History of the Tesseract OCR engine: what worked and what didn't. In R. Zanibbi & B. Coüasnon (Eds.), *Document recognition and retrieval xx* (p. 865802). SPIE. https://doi.org/10.1117/12.2010051

Wikipedia:popular pages. (2024). https://en.wikipedia.org/wiki/Wikipedia:Popular_pages

Appendix A

Wikipedia Article Dataset

Category	Articles
People	Donald Trump, Elizabeth II, Barack Obama, Christiano Ronaldo, Michael Jackson, Elon Musk, Lady Gaga, Adolf Hitler, Eminem, Lionel Messi, Justin Bieber, Freddie Mercury, Kim Kardashian, Johnny Depp, Steve Jobs, Dwayne Johnson, Michael Jordan, Taylor Swift, Stephen Hawking, Kanye West
Present countries	United States, India, United Kingdom, Canada, Australia, China, Russia, Japan, Germany, France, Singapore, Israel, Pakistan, Philippines, Brazil, Italy, Netherlands, New Zealand, Ukraine, Spain
Cities	New York City, London, Hong Kong, Los Angeles, Dubai, Washington, D.C., Paris, Chicago, Angelsberg, Mumbai, San Francisco, Rome, Monaco, Toronto, Tokyo, Philadelphia, Machu Picchu, Jerusalem, Amsterdam, Boston
Life	Cat, Dog, Animal, Lion, Coronavirus, Tiger, Human, Dinosaur, Elephant, Virus, Horse, Photosynthesis, Evolution, Apple, Bird, Mammal, Potato, Polar bear, Shark, Snake
Buildings and structures	Taj Mahal, Burj Khalifa, Statue of Liberty, Great Wall of China, Eiffel Tower, Berlin Wall, Stonehenge, Mount Rushmore, Colosseum, Auschwitz concentration camp, Great Pyramid of Giza, One World Trade Center, Empire State Building, White House, Petra, Large Hadron Collider, Hagia Sophia, Golden Gate Bridge, Panama Canal, Angkor Wat

Table A.1: Dataset of 100 Wikipedia articles